







a,b: $n = 100$ independent development scenarios per candidate; finite grid. c: $n = 300$; failure families may overlap.
d: separate 986000-series count test; both service scores coincide at each request; no binomial confidence inference.
Matched last-call outcomes remain in Table 15; historical selection and score interpretation are in Supplementary Note 7.